

[Get started](#)[Open in app](#)

towards
data science

[Follow](#)

580K Followers



ID3 Decision Tree Classifier from scratch in Python

Coding the ID3 algorithm to build a Decision Tree Classifier from scratch.



[Bernardo Garcia del Rio](#) Dec 13, 2020 · 7 min read



Photo by [Fabrice Villard](#) on [Unsplash](#)

In my last article, I showed how you could have some fun **growing a Random Forest using Sklearn's DecisionTreeClassifier**. You can check it out below:

[Get started](#)[Open in app](#)

Understanding Decision Trees and Random Forests with a hands-on example

bernardogarciadelrio.medium.com



If you are new to Machine Learning (ML) or are in the process of becoming a Data Scientist or ML practitioner like me, you will probably get something out of it. In my case, preparing the material for the article helped me gain a good understanding of how Random Forests and Decision Trees work under the hood without getting into complex code.


In that article, I mentioned that there are many algorithms that can be used to build a Decision Tree. One of them is **ID3 (Iterative Dichotomiser 3)** and we are going to see how to code it from scratch using ONLY Python to build a Decision Tree Classifier.

All the code can be found in a public repository that I have attached below:

bergr7/ID3_From_Scratch

All the code is in a public repository at the link below: Installation Project Motivation File Descriptions Results...

[github.com](https://github.com/bergr7/ID3_From_Scratch)



Modelling the nodes of the tree

A Decision Tree is formed by nodes: root node, internal nodes and leaf nodes. We can create a Python class that will contain all the information of all the nodes of the Decision Tree.

```
1 class Node:
2     """Contains the information of the node and another nodes of the Decision Tree."""
3
4     def __init__(self):
5         self.value = None
6         self.next = None
7         self.childs = None
8
```

nodes.py hosted with ❤ by GitHub

[view raw](#)

[Get started](#)[Open in app](#)

- *value*: Feature to make the split and branches.
- *next*: Next node
- *childs*: Branches coming off the decision nodes

Decision Tree Classifier Class

We create now our main class called `DecisionTreeClassifier` and use the `__init__` constructor to initialise the attributes of the class and some important variables that are going to be needed.

Note that I have provided many annotations in the code snippets that help understand the code.

```

1  class DecisionTreeClassifier:
2      """Decision Tree Classifier using ID3 algorithm."""
3
4      def __init__(self, X, feature_names, labels):
5          self.X = X # features or predictors
6          self.feature_names = feature_names # name of the features
7          self.labels = labels # categories
8          self.labelCategories = list(set(labels)) # unique categories
9          # number of instances of each category
10         self.labelCategoriesCount = [list(labels).count(x) for x in self.labelCategories]
11         self.node = None # nodes
12         # calculate the initial entropy of the system
13         self.entropy = self._get_entropy([x for x in range(len(self.labels))])
14

```

decisiontree.py hosted with ❤ by GitHub

[view raw](#)

In order to calculate the entropy note we are using the private function `_get_entropy()`. The code for this function is provided below:

```

1  def _get_entropy(self, x_ids):
2      """ Calculates the entropy.
3      Parameters
4      _____
5      :param x_ids: list, List containing the instances ID's
6      _____
7      :return: entropy: float, Entropy.
8      """
9      # sorted labels by instance id
10     labels = [self.labels[i] for i in x_ids]

```

[Get started](#)[Open in app](#)

```
14     entropy = sum([-count / len(x_ids) * math.log(count / len(x_ids), 2)
15                   if count else 0
16                   for count in label_count
17                   ])
18
19     return entropy
20
```

entropy.py hosted with ❤ by GitHub

[view raw](#)

We pass the instances id's or indexes to this function. For doing this, we need to generate an unique number for each instance. Python's lists comprehensions come in very handy for this task as you can see.

We are going to code an ID3 algorithm that uses the **information gain** to find the feature that maximises it and make a split based on that feature. The information gain is based on entropy.

As this article is about coding the ID3 algorithm, I am not going to go into the details but **entropy** is a measure of uncertainty about a random variable. If we minimise the entropy, then we increase the certainty about the variable. In other words, if the random variable can take only one value the entropy reaches its minimum whereas if all the values are equiprobable the entropy is maximum.

The algorithm will try to minimise the entropy or, equivalently, maximising the information gain.

We can compute the entropy with the following formula:

$$H(S) = \sum_{i=1}^N -p_i \log_2 p_i$$

Entropy formula (Caption by Author)

where p_i is the proportion of each category $i=1,\dots,N$.

We said that we would compute the information gain to choose the feature that maximises it and then make the split based on that feature. The information gain is a

[Get started](#)[Open in app](#)

$$G(S, j) = H(S) - \sum_j \frac{|S_j|}{|S|} H(S_j)$$

Information gain (Caption by Author)

So we create another private function that computes the information gain:

```

1  def _get_information_gain(self, x_ids, feature_id):
2      """Calculates the information gain for a given feature based on its entropy and the
3      Parameters
4      _____
5      :param x_ids: list, List containing the instances ID's
6      :param feature_id: int, feature ID
7      _____
8      :return: info_gain: float, the information gain for a given feature.
9      """
10     # calculate total entropy
11     info_gain = self._get_entropy(x_ids)
12     # store in a list all the values of the chosen feature
13     x_features = [self.X[x][feature_id] for x in x_ids]
14     # get unique values
15     feature_vals = list(set(x_features))
16     # get frequency of each value
17     feature_v_count = [x_features.count(x) for x in feature_vals]
18     # get the feature values ids
19     feature_v_id = [
20         [x_ids[i]
21          for i, x in enumerate(x_features)
22          if x == y]
23         for y in feature_vals
24     ]
25
26     # compute the information gain with the chosen feature
27     info_gain_feature = sum([v_counts / len(x_ids) * self._get_entropy(v_ids)
28                             for v_counts, v_ids in zip(feature_v_count, feature_v_id)])
29
30     info_gain = info_gain - info_gain_feature
31
32     return info_gain
33

```

infogain.py hosted with ❤ by GitHub

[view raw](#)

[Get started](#)[Open in app](#)

the feature specified in `feature_id` and finally the information gain.

This ID3 algorithm chooses the feature that maximise the information gain at each split. That is if the amount information in the feature is large, the entropy will be small, and therefore the second term of the information gain formula will be also small, increasing the information gain.

`_get_information_gain()` only calculates the information gain for a given set and feature but, at each split, we need to figure out which is the **feature** that maximises the information gain. We need to create another function to find the feature that maximises the information gain.

```

1  def _get_feature_max_information_gain(self, x_ids, feature_ids):
2      """Finds the attribute/feature that maximizes the information gain.
3      Parameters
4      _____
5      :param x_ids: list, List containing the samples ID's
6      :param feature_ids: list, List containing the feature ID's
7      _____
8      :returns: string and int, feature and feature id of the feature that maximizes the i
9      """
10     # get the entropy for each feature
11     features_entropy = [self._get_information_gain(x_ids, feature_id) for feature_id in
12     # find the feature that maximises the information gain
13     max_id = feature_ids[features_entropy.index(max(features_entropy))]
14
15     return self.feature_names[max_id], max_id

```

getmaxinfofain.py hosted with ❤ by GitHub

[view raw](#)

Ok..so we have a private function (`_get_feature_max_information_gain()`) that find the feature that maximises the information gain, that uses another private function (`_get_information_gain()`) that calculates the information gain, that uses another private function (`_get_entropy()`) that computes the entropy. We are all set! We can start coding the ID3 algorithm that will create our ID3 Decision Tree for classification problems.

We create a function that initialises the algorithm and then uses a private function to call the algorithm recursively to build our tree.

```

1  def id3(self):
2      """Initializes ID3 algorithm to build a Decision Tree Classifier.

```

[Get started](#)[Open in app](#)

```
6      # assign an unique number to each instance
7      x_ids = [x for x in range(len(self.X))]
8      # assign an unique number to each feature
9      feature_ids = [x for x in range(len(self.feature_names))]
10     # define node variable - instance of the class Node
11     self.node = self._id3_recv(x_ids, feature_ids, self.node)
```

id3.py hosted with ❤ by GitHub

[view raw](#)

Note how we call `_id3_recv()` function in `self.node`. This function returns an instance of the class `Node` with information of all the nodes (decisions) in the Decision Tree.

`_id3_recv()` is the trickiest function to code so let's spend some time understanding what it does.

Let's generate some data

I have generated a small toy dataset so we can work on a small example and actually visualise the resulting tree. This will help us understand the algorithm as well.

Let's imagine that we love surfing waves and would like to have a model to predict if there will be good waves in our favourite spot.

Our features are:

Feature	Values
wind_direction	{ 'N', 'S', 'E', 'W' }
tide	{ 'high', 'low' }
swell_forecasting	{ 'small', 'medium', 'large' }

Features (Caption by Author)

And our target is simply a binary categorical variable that tells us if there were/will be good waves based on the information provided in our features:

Target	Values
good_waves	{ 'Yes', 'No' }

Target (Caption by Author)

Let's generate some synthetic data using random sampling:

[Get started](#)[Open in app](#)

```

4     'tide': ['Low', 'High'],
5     'swell_forecasting': ['small', 'medium', 'large'],
6     'good_waves': ['Yes', 'No']
7 }
8
9 # create an empty dataframe
10 data_df = pd.DataFrame(columns=data.keys())
11
12 np.random.seed(42)
13 # randomly create 1000 instances
14 for i in range(1000):
15     data_df.loc[i, 'wind_direction'] = str(np.random.choice(data['wind_direction'], 1)[0])
16     data_df.loc[i, 'tide'] = str(np.random.choice(data['tide'], 1)[0])
17     data_df.loc[i, 'swell_forecasting'] = str(np.random.choice(data['swell_forecasting'], 1)[0])
18     data_df.loc[i, 'good_waves'] = str(np.random.choice(data['good_waves'], 1)[0])
19
20 data_df.head()

```

data.py hosted with ❤ by GitHub

[view raw](#)

	wind_direction	tide	swell_forecasting	good_waves
0	E	Low	large	No
1	E	Low	large	No
2	W	Low	small	No
3	N	High	small	No
4	N	High	small	Yes

First 5 rows of our toy dataset (Caption by Author)

`_id3_recv()` function

Now that we have some data to work with it will be easier to understand `_id3_recv()` function. Below is the code with some annotations:

```

1 def _id3_recv(self, x_ids, feature_ids, node):
2     """ID3 algorithm. It is called recursively until some criteria is met.
3     Parameters
4     _____
5     :param x_ids: list, list containing the samples ID's
6     :param feature_ids: list, List containing the feature ID's
7     :param node: object, An instance of the class Nodes
8     _____
9     :returns: An instance of the class Node containing all the information of the nodes
10    """
11    if not node:

```


[Get started](#)[Open in app](#)

```

15     # if all the example have the same class (pure node), return node
16     if len(set(labels_in_features)) == 1:
17         node.value = self.labels[x_ids[0]]
18         return node
19     # if there are not more feature to compute, return node with the most probable class
20     if len(feature_ids) == 0:
21         node.value = max(set(labels_in_features), key=labels_in_features.count) # compute
22         return node
23     # else...
24     # choose the feature that maximizes the information gain
25     best_feature_name, best_feature_id = self._get_feature_max_information_gain(x_ids, f
26     node.value = best_feature_name
27     node.childs = []
28     # value of the chosen feature for each instance
29     feature_values = list(set([self.X[x][best_feature_id] for x in x_ids]))
30     # loop through all the values
31     for value in feature_values:
32         child = Node()
33         child.value = value # add a branch from the node to each feature value in our f
34         node.childs.append(child) # append new child node to current node
35         child_x_ids = [x for x in x_ids if self.X[x][best_feature_id] == value]
36         if not child_x_ids:
37             child.next = max(set(labels_in_features), key=labels_in_features.count)
38             print('')
39         else:
40             if feature_ids and best_feature_id in feature_ids:
41                 to_remove = feature_ids.index(best_feature_id)
42                 feature_ids.pop(to_remove)
43                 # recursively call the algorithm
44                 child.next = self._id3_recv(child_x_ids, feature_ids, child.next)
45     return node

```

id3recv.py hosted with ❤ by GitHub

[view raw](#)

The function takes two lists containing instances ids and feature ids and an instance of the class Node. Firstly, it checks if the node attribute is *None* and if it is, it initialises an instance of the class Node.

Then, at each split the algorithm checks if some of the following conditions are met:

- **Are all the instances of the same category?** If so, it assigns the category to the node value and returns the variable node. The following chunk of code does this task in the function:

```

1  # if all the example have the same class (pure node), return node

```

[Get started](#)[Open in app](#)

condition1.py hosted with ❤ by GitHub

[view raw](#)

- **Are there more features to compute the information gain?** If there are not, it assigns the most probably category to the node value (it computes the mode) and returns the variable node. Code doing this is below:

```

1  # if there are not more feature to compute, return node with the most probable class
2      if len(feature_ids) == 0:
3          node.value = max(set(labels_in_features), key=labels_in_features.count) # co
4      return node

```

condition2.py hosted with ❤ by GitHub

[view raw](#)

If none of the above conditions are met, it chooses the feature that maximises the information gain and stores it in the node value. In this case, the node value is the **feature** with which it is making the split.

The ID3 algorithm creates a branch for each value of the selected feature and finds the instances in the training set that takes that branch. Note each branch is represented with a new instance of the class node that also contains the the next node.

For the next node, it computes the information gain for the remaining features and instances and chooses the one that maximises it to make another split until all the instances have the same class or there not more features to compute, returning the most probable category in this case. All of this is done in the last part of the function:

```

1      # else...
2      # choose the feature that maximizes the information gain
3      best_feature_name, best_feature_id = self._get_feature_max_information_gain(x_ids,
4      node.value = best_feature_name
5      node.childs = []
6      # value of the chosen feature for each instance
7      feature_values = list(set([self.X[x][best_feature_id] for x in x_ids]))
8      # loop through all the values
9      for value in feature_values:
10         child = Node()
11         child.value = value # add a branch from the node to each feature value in our
12         node.childs.append(child) # append new child node to current node
13         child_x_ids = [x for x in x_ids if self.X[x][best_feature_id] == value] # insta
14         if not child_x_ids:
15             child.next = max(set(labels_in_features), key=labels_in_features.count)

```

[Get started](#)[Open in app](#)

```

19         to_remove = feature_ids.index(best_feature_id)
20         feature_ids.pop(to_remove)
21         # recursively call the algorithm
22         child.next = self._id3_recv(child_x_ids, feature_ids, child.next)
23     return node

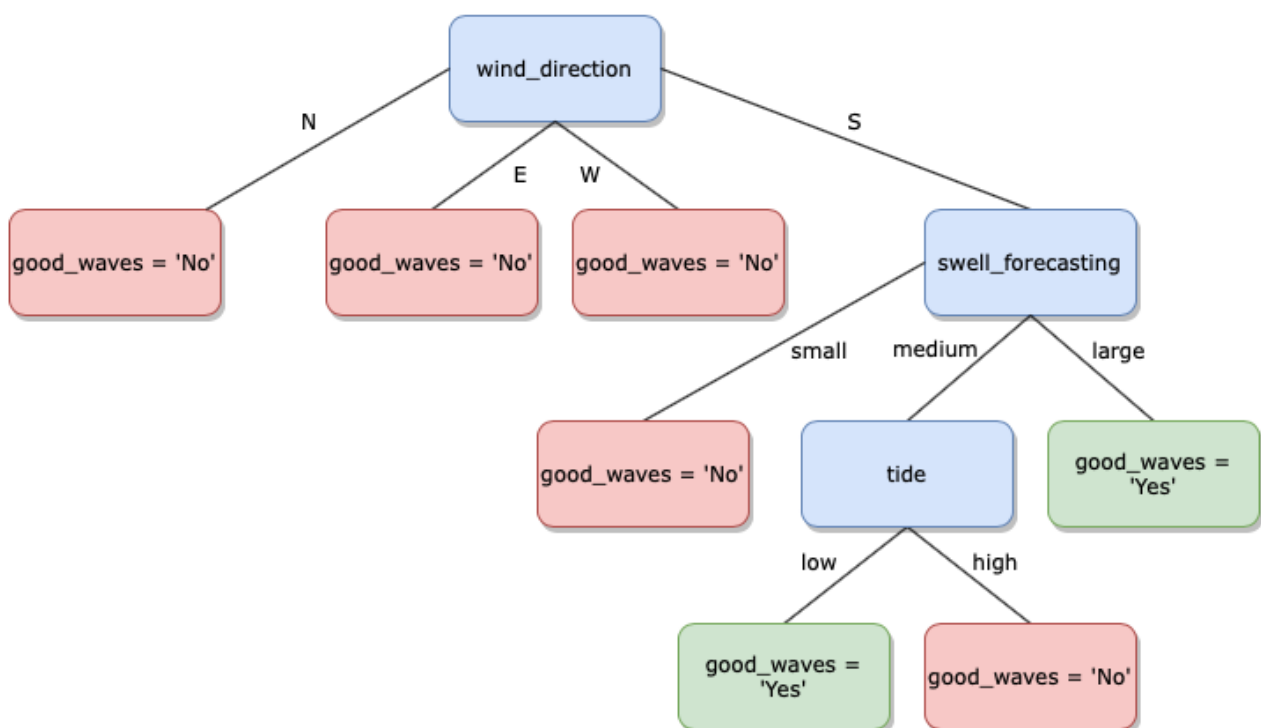
```

id3process.py hosted with ❤ by GitHub

[view raw](#)

Resulting Decision Tree

And here we have the Decision Tree Classifier that the ID3 algorithm has built using our synthetic data:



Resulting Decision Tree (Caption by Author)

It seems that our “synthetic” beach only works well when the wind blows from the South. If there is a large swell, it is likely that there will be good waves regardless of the tide but if the swell is medium, we will usually find better surfing conditions during low tide.

I would like to give credit to the Professors of the MSc in Machine Learning, Deep Learning and AI by Big Data International Campus for providing a version of the source code, that inspired me to write this article.

[Get started](#)[Open in app](#)

International Campus and UCAM

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)



Get this newsletter

[Decision Tree](#)[Decision Tree Classifier](#)[Machine Learning](#)[Id3 Algorithm](#)[Artificial Intelligence](#)[About](#) [Write](#) [Help](#) [Legal](#)

Get the Medium app



Download on the
App Store



GET IT ON
Google Play